

ONLINE LEARNING APPLICATIONS

Project: Dynamic Pricing with Production Constraints

PRESENTED BY:

Simar Achmet Kechagia

Berke Chalishkan

Alperen Dagi

Nikola Dimic

Veljko Tatalovic

Requirements

01	Single product and stochastic environment
02	Multiple products and stochastic environment
03	Best-of-both-worlds algorithms with a single product
04	Best-of-both-worlds with multiple products
05	Slightly non-stationary environments with multiple products

Requirement 1 : Single product and stochastic environment

Environment:

- Valuation distribution sampled from different Beta distributions.

Algorithms implemented:

- UCB1 (ignoring inventory)
- UCB1 with Inventory Constraint

Evaluation:

- Cumulative Reward
 - Cumulative Regret vs \sqrt{T} and T bounds
-

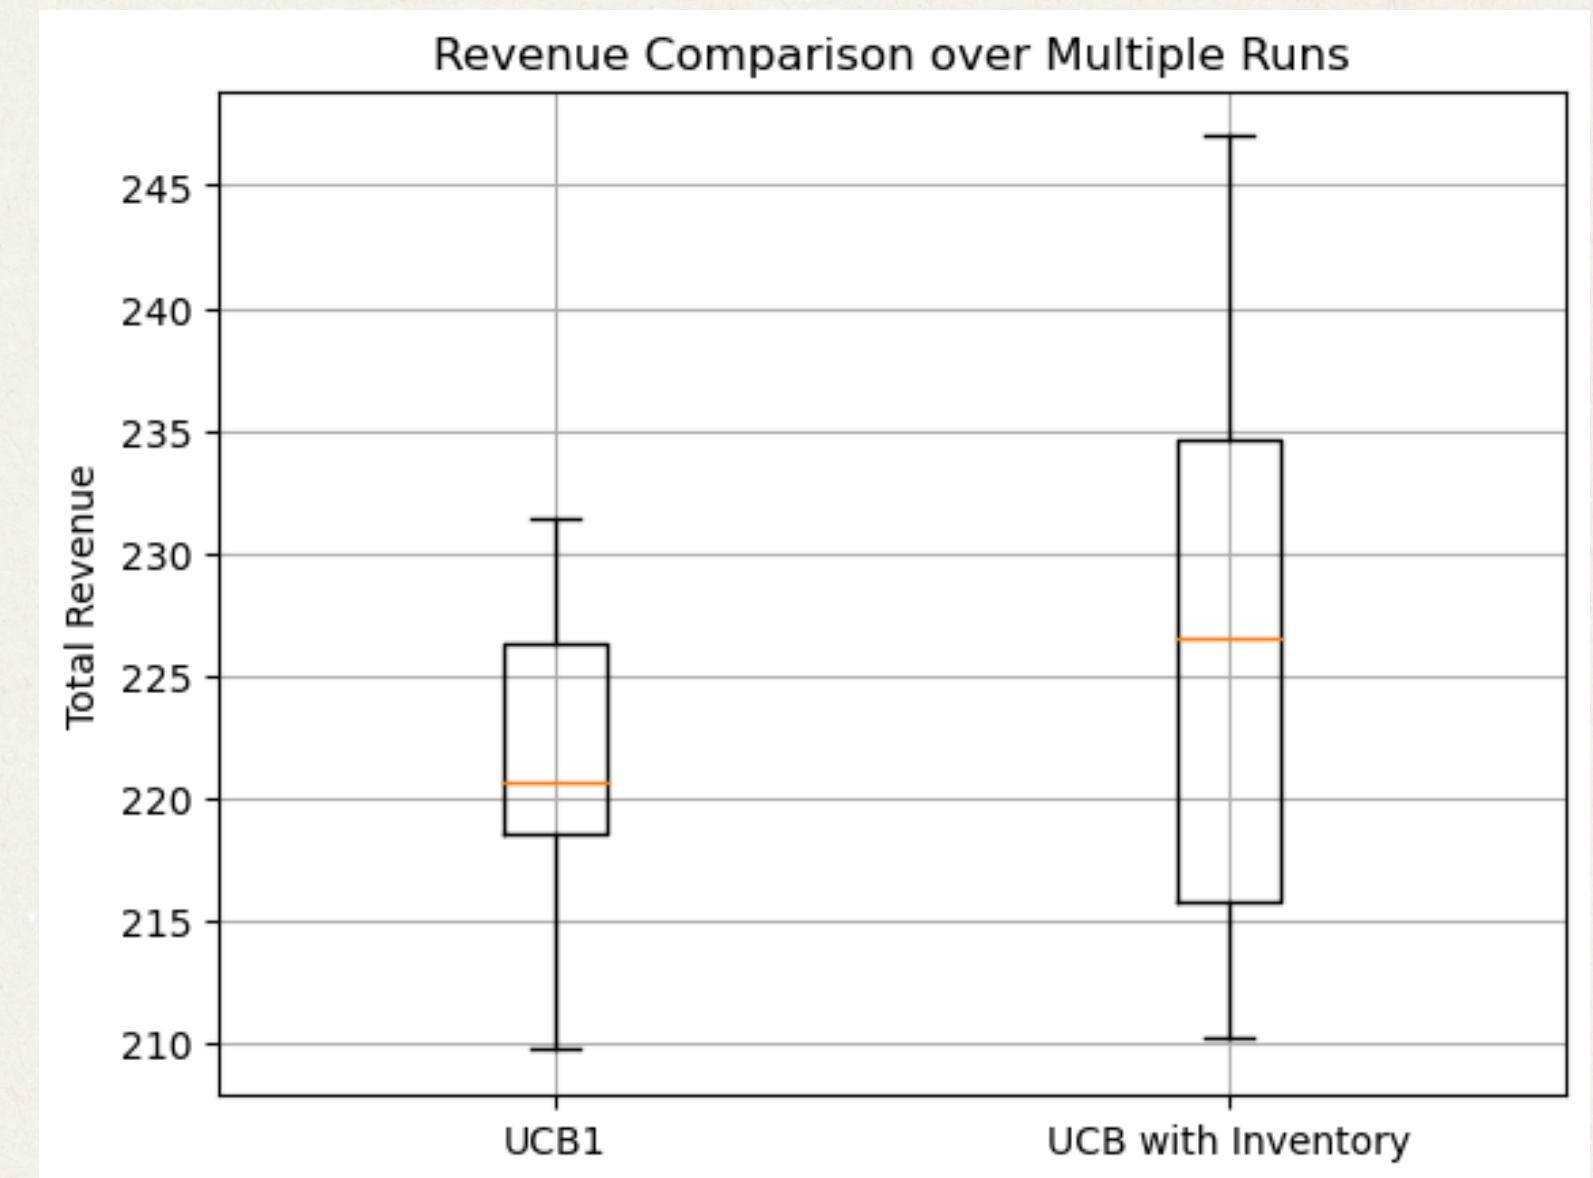
Requirement 1 : Single product and stochastic environment

Observations:

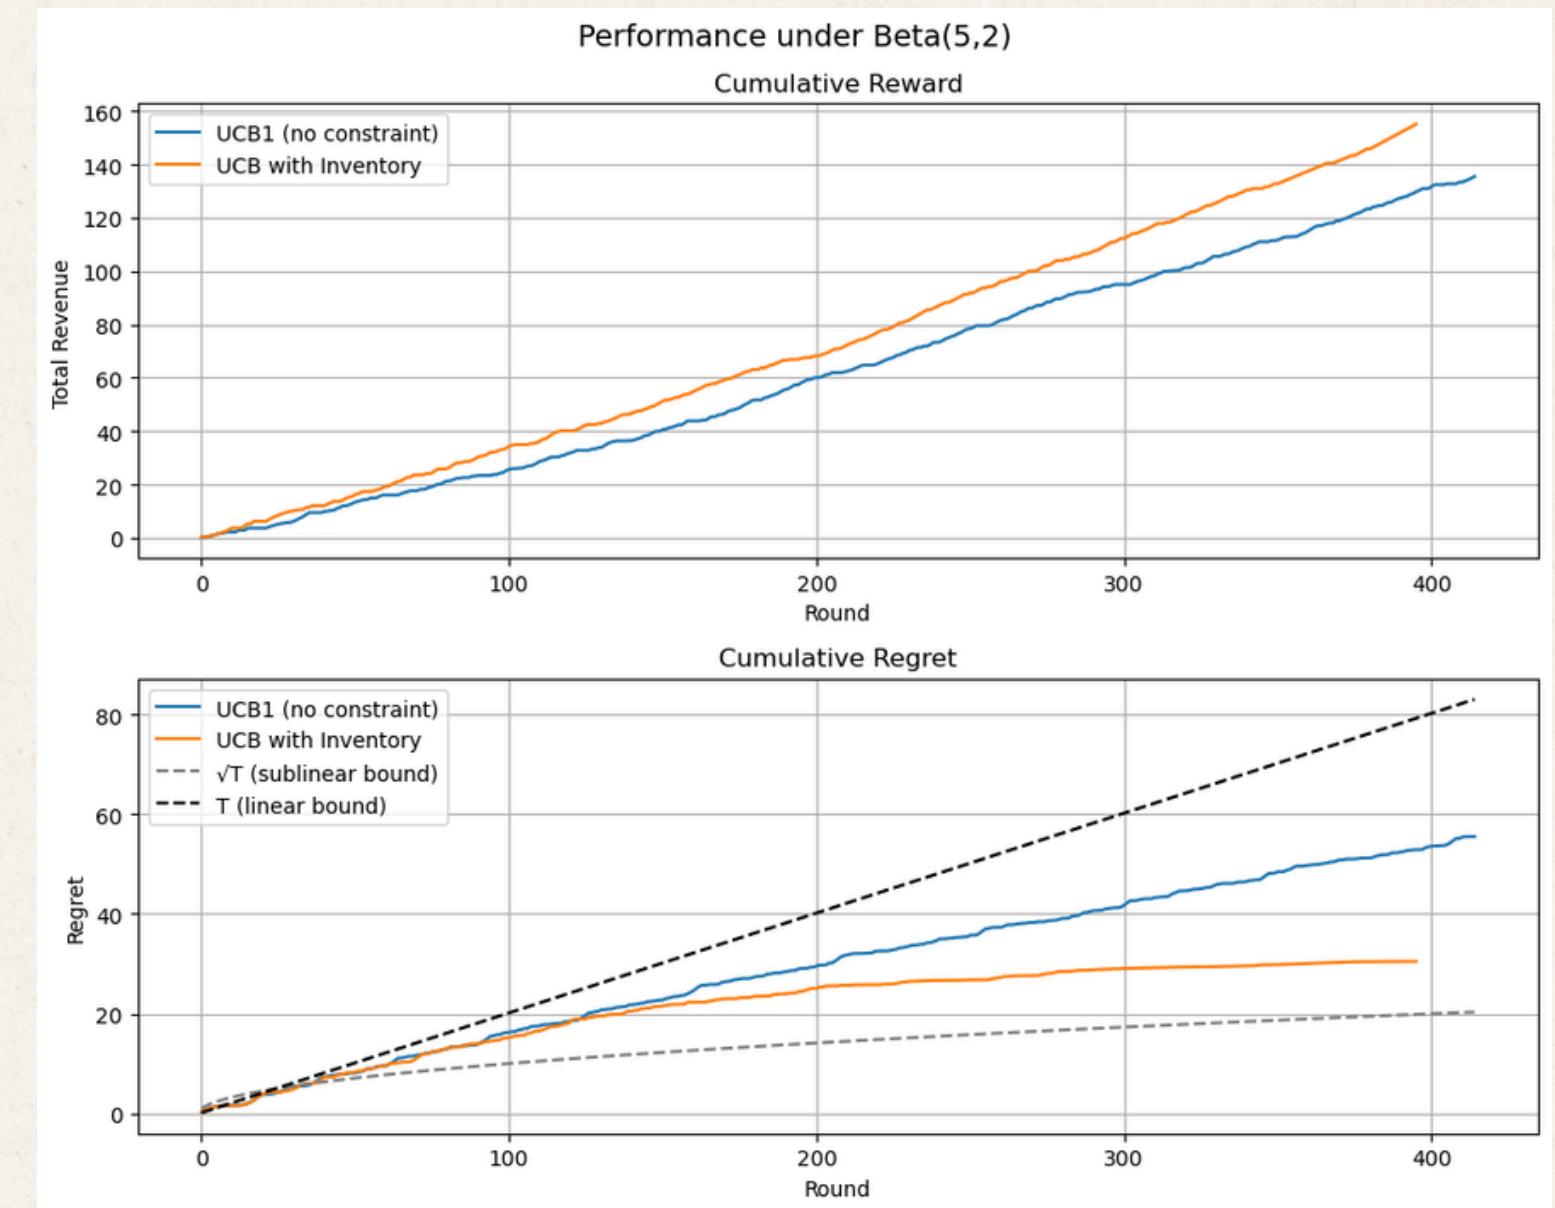
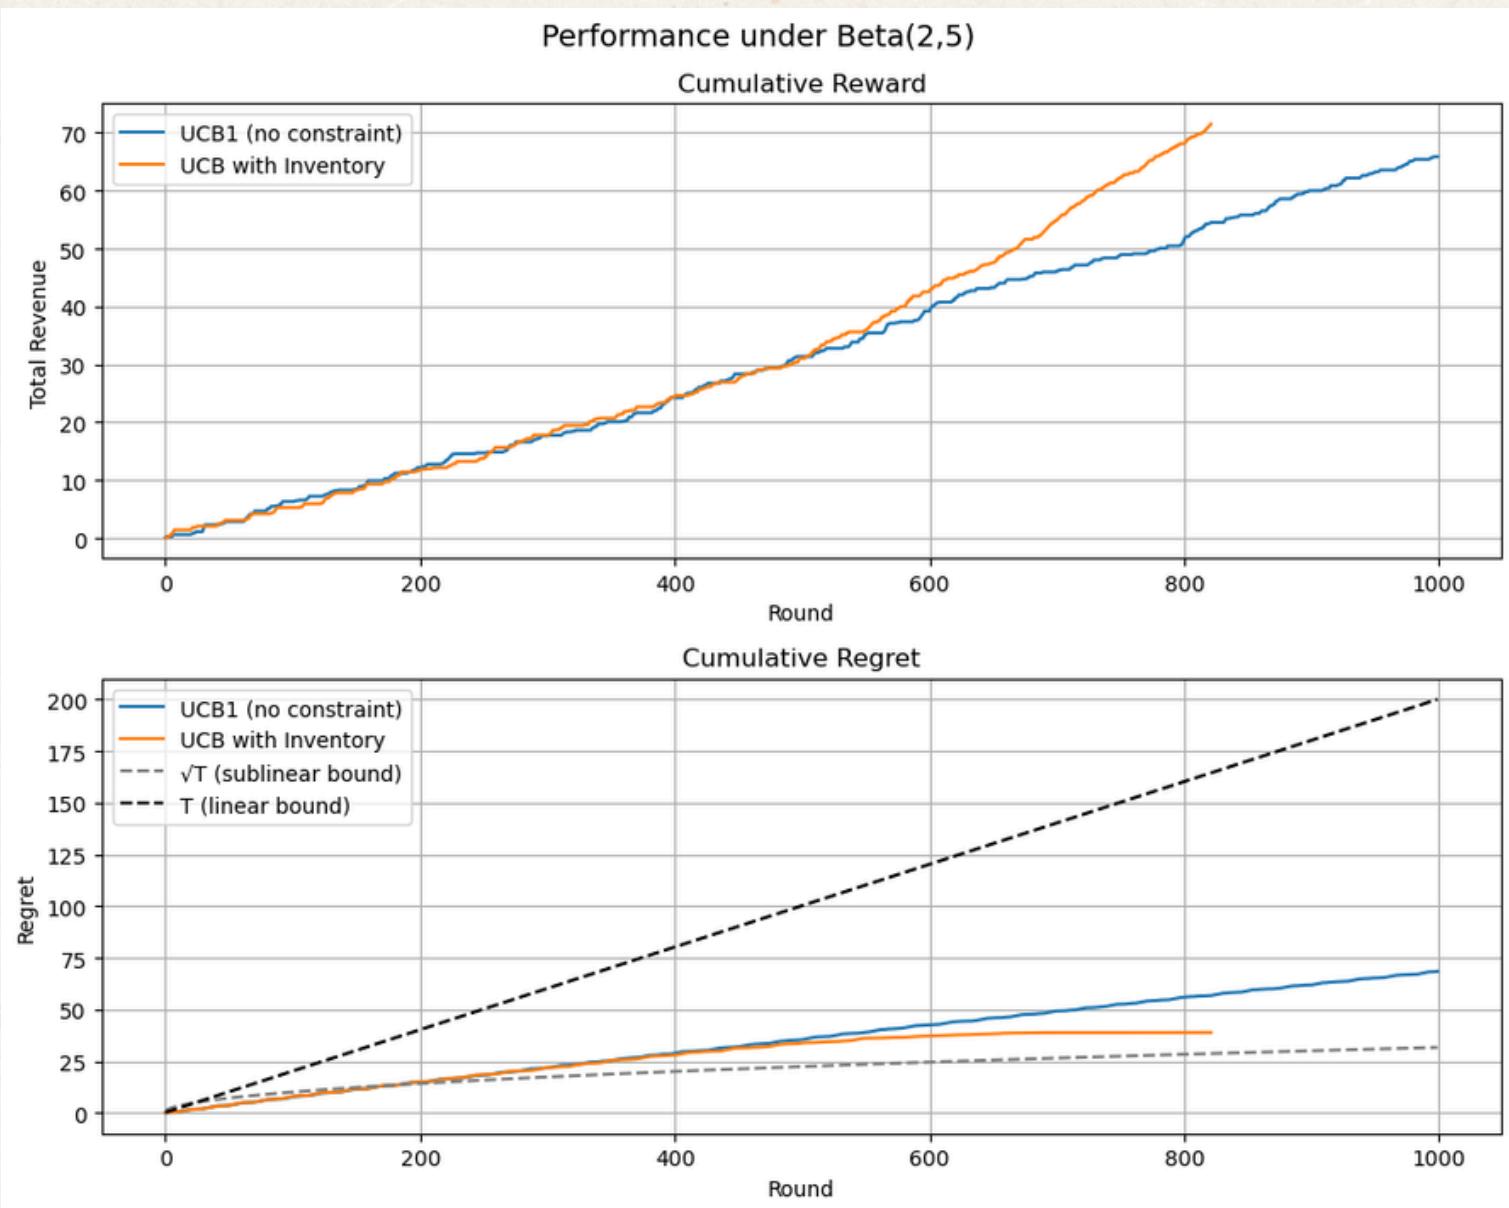
- UCB1 sometimes outperforms when the budget isn't binding.
- Inventory-aware UCB1 is safer and more robust when the budget is tight.

Settings:

- Budget **B = 1000**
- Horizon **T = 10000**
- Prices: **{0.1, 0.2, 0.3, 0.4, 0.5, 0.6, 0.7, 0.8, 0.9, 1.0}**
- Simulations: **20**



Requirement 1 : Single product and stochastic environment



Requirement 2 : Multiple products and stochastic environment

Environment:

- Heterogeneous valuation distributions per product.

Algorithms implemented:

- Combinatorial-UCB with inventory constraint.

Evaluation:

- Cumulative Reward
- Cumulative Regret vs Oracle
- Price Histograms

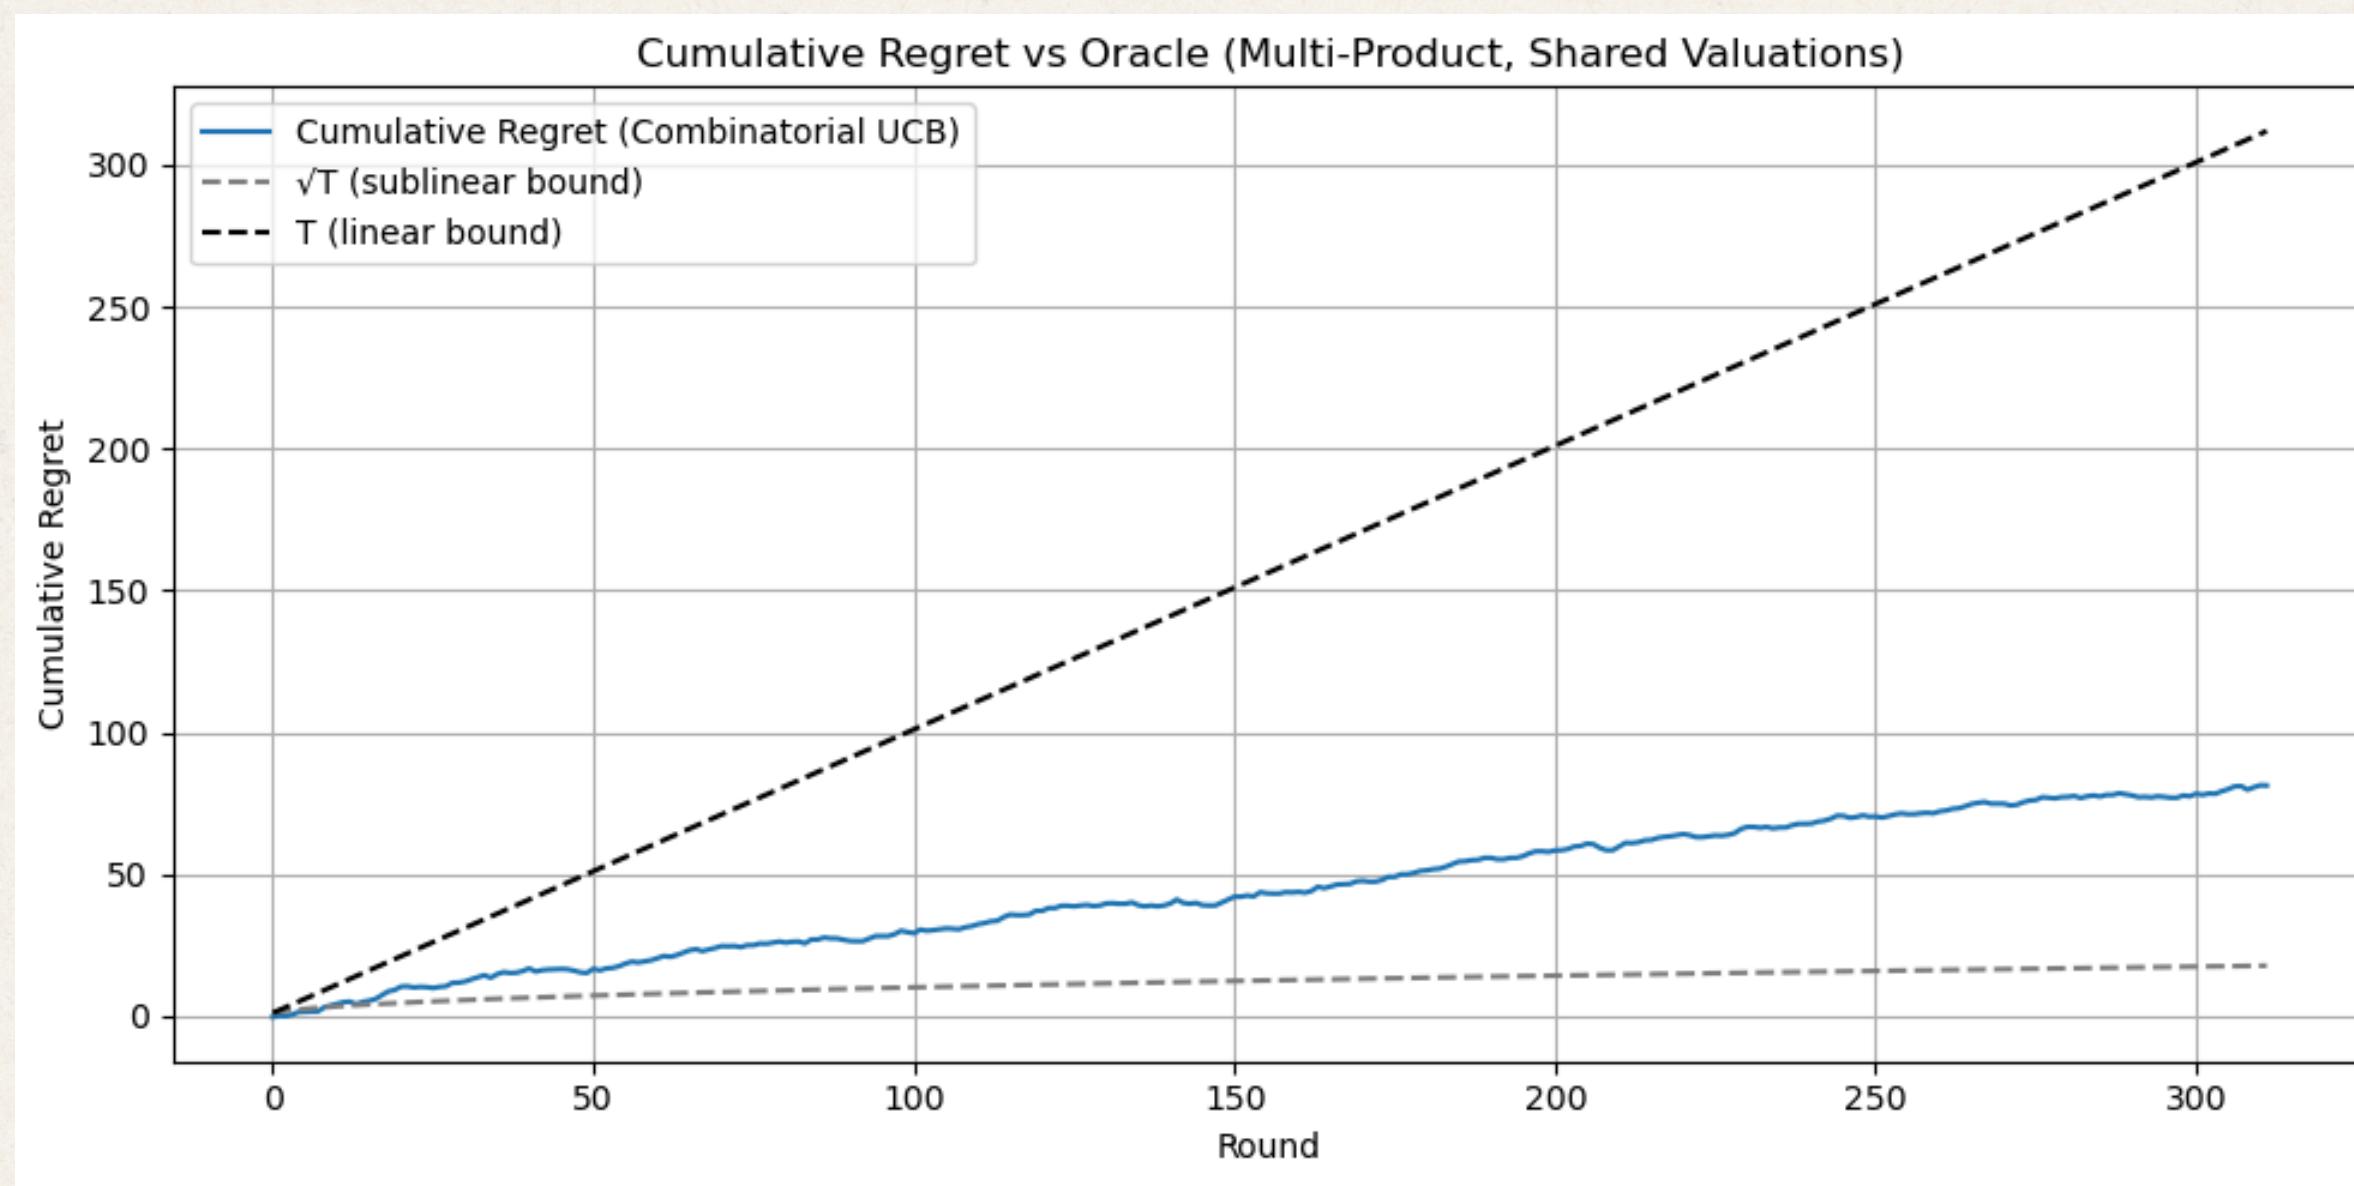
Observations:

- Combinatorial-UCB adapts independently per product.
- Doesn't outperform the Greedy baseline in total revenue. ***

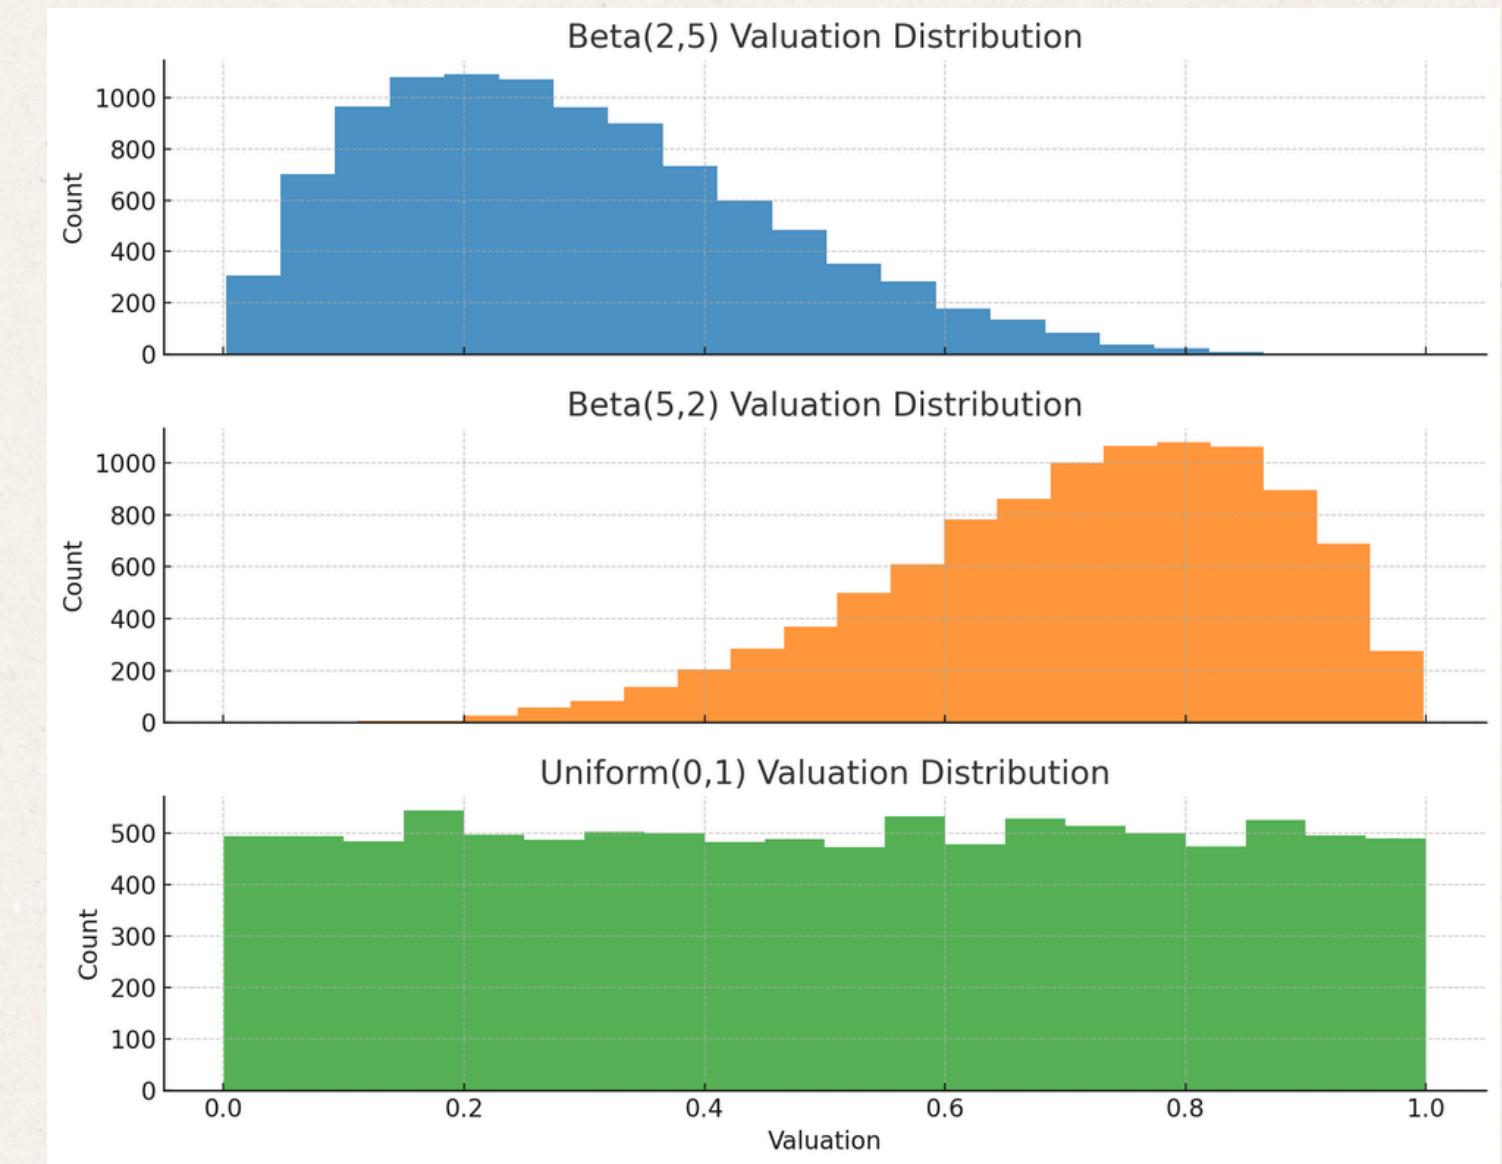
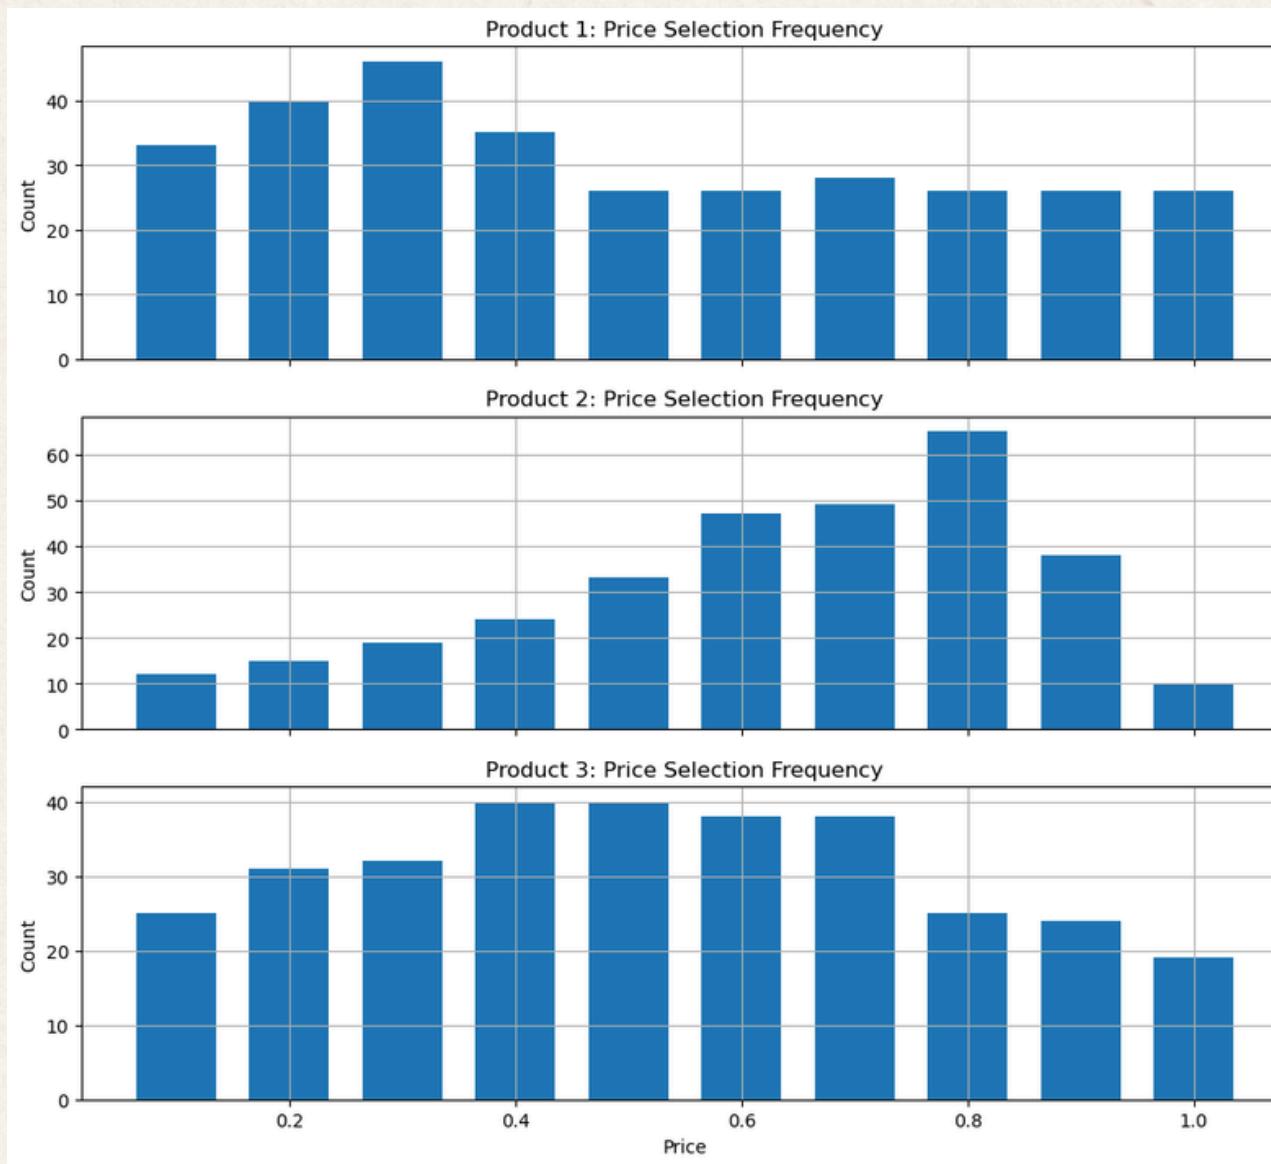
Settings:

- Budget **B = 500**
- Horizon **T = 1000**
- Prices: **{0.1, 0.2, 0.3, 0.4, 0.5, 0.6, 0.7, 0.8, 0.9, 1.0}**

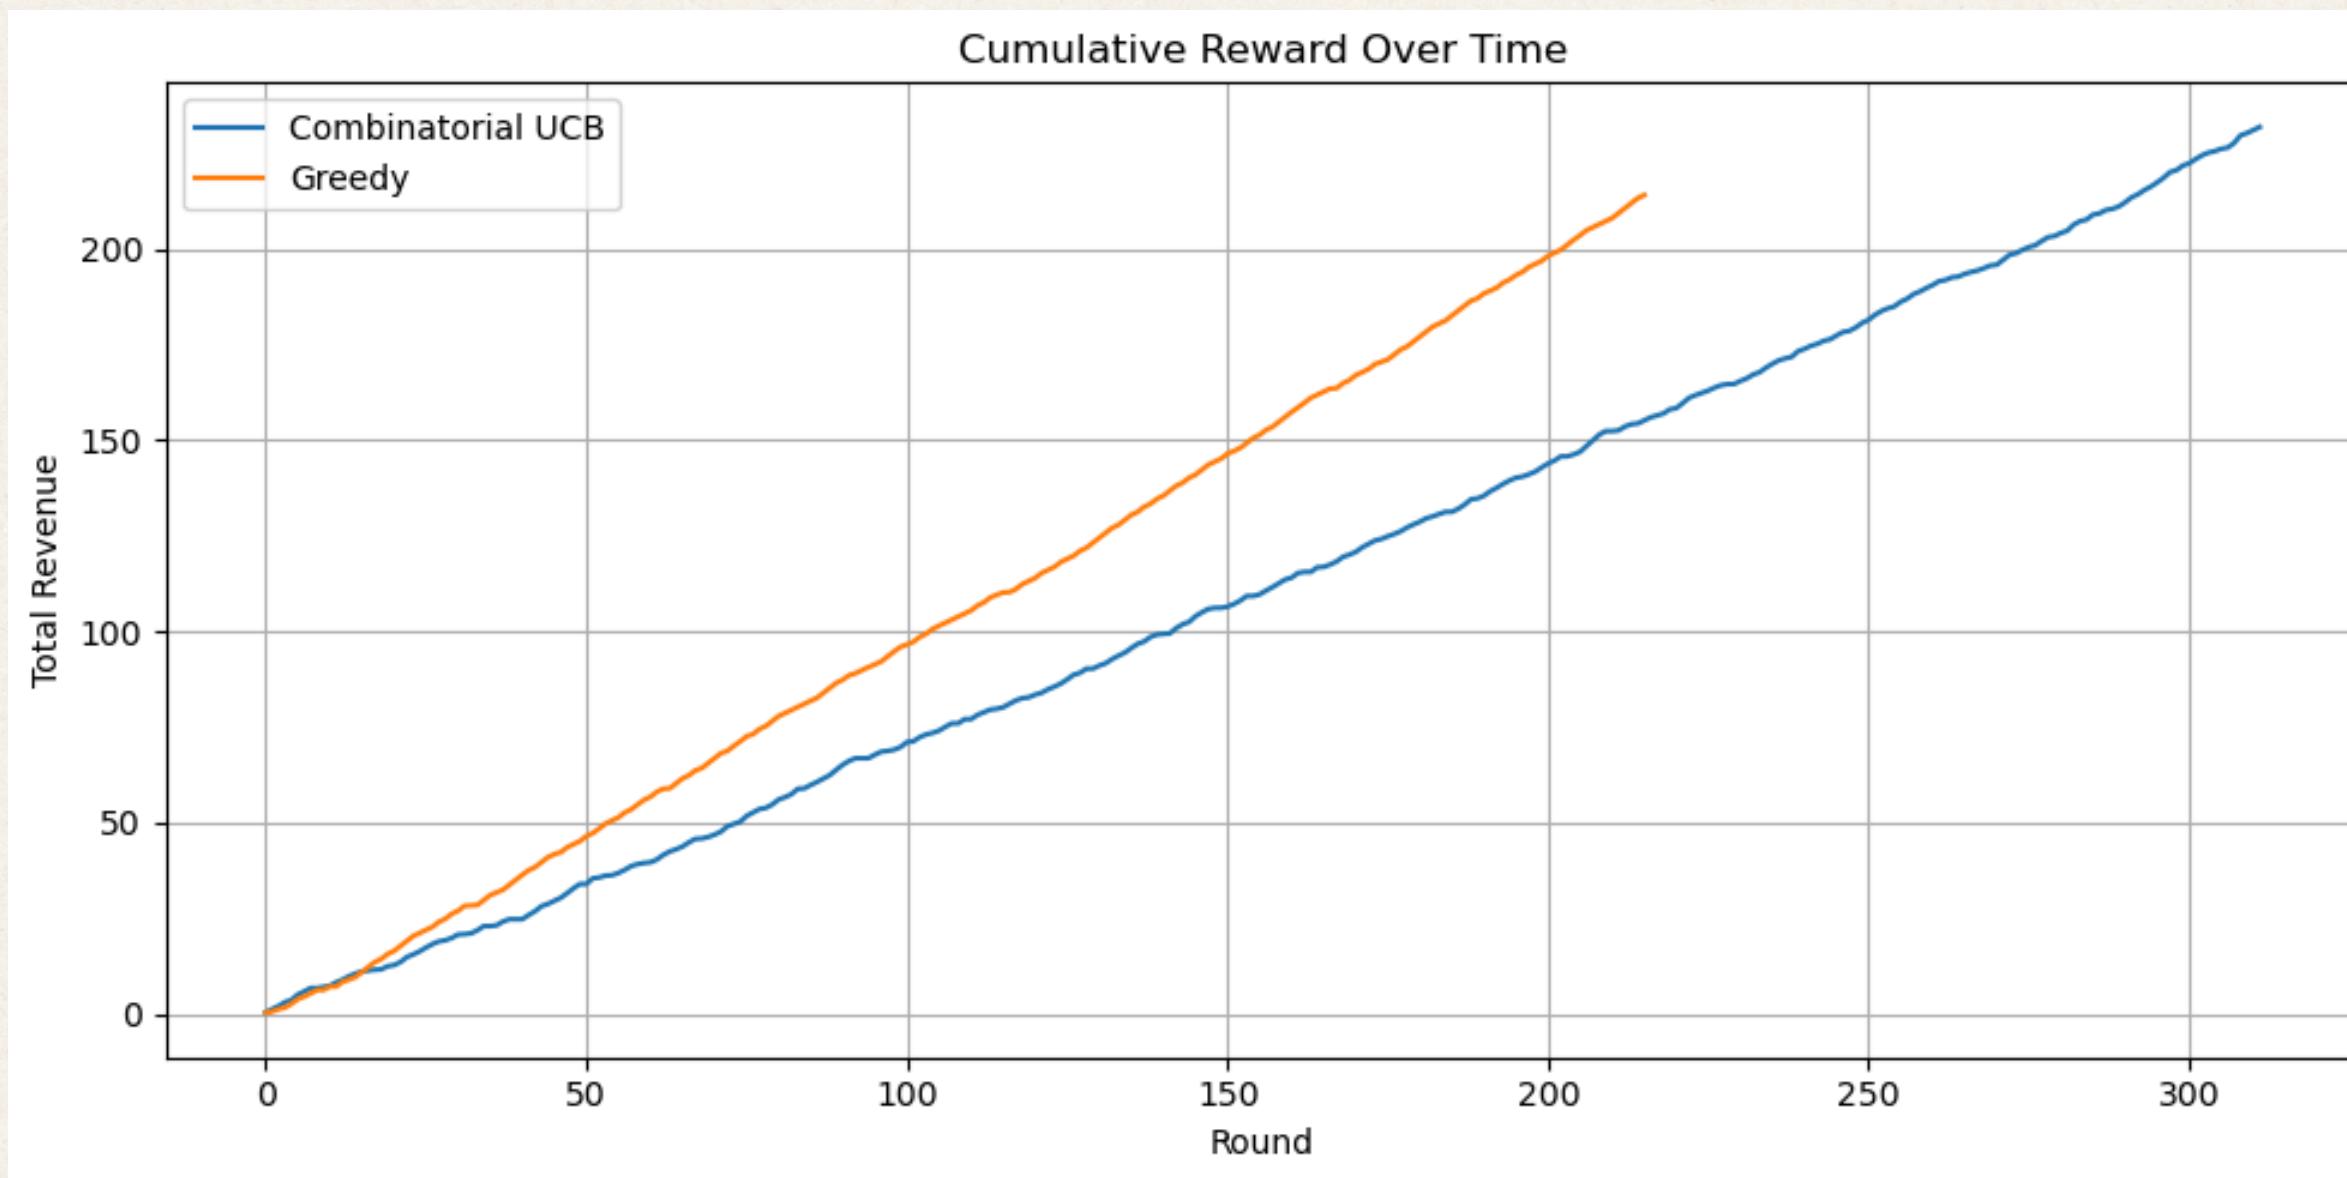
Requirement 2 : Multiple products and stochastic environment



Requirement 2 : Multiple products and stochastic environment



Requirement 2 : Multiple products and stochastic environment



Requirement 3: Best-of-Both-Worlds and single product

Environment:

- Highly non-stationary valuation distribution.

Algorithms implemented:

- Primal-Dual pricing strategy.

Method:

- Dual variable λ updated as:

$$\lambda_{t+1} = \left[\lambda_t + \eta \cdot \left(x_t - \frac{B}{T} \right) \right]^+$$

- Prices chosen as:

$$p_t = \arg \max_{p \in P} (p - \lambda_t) \cdot \mathbb{I}(v_t \geq p)$$

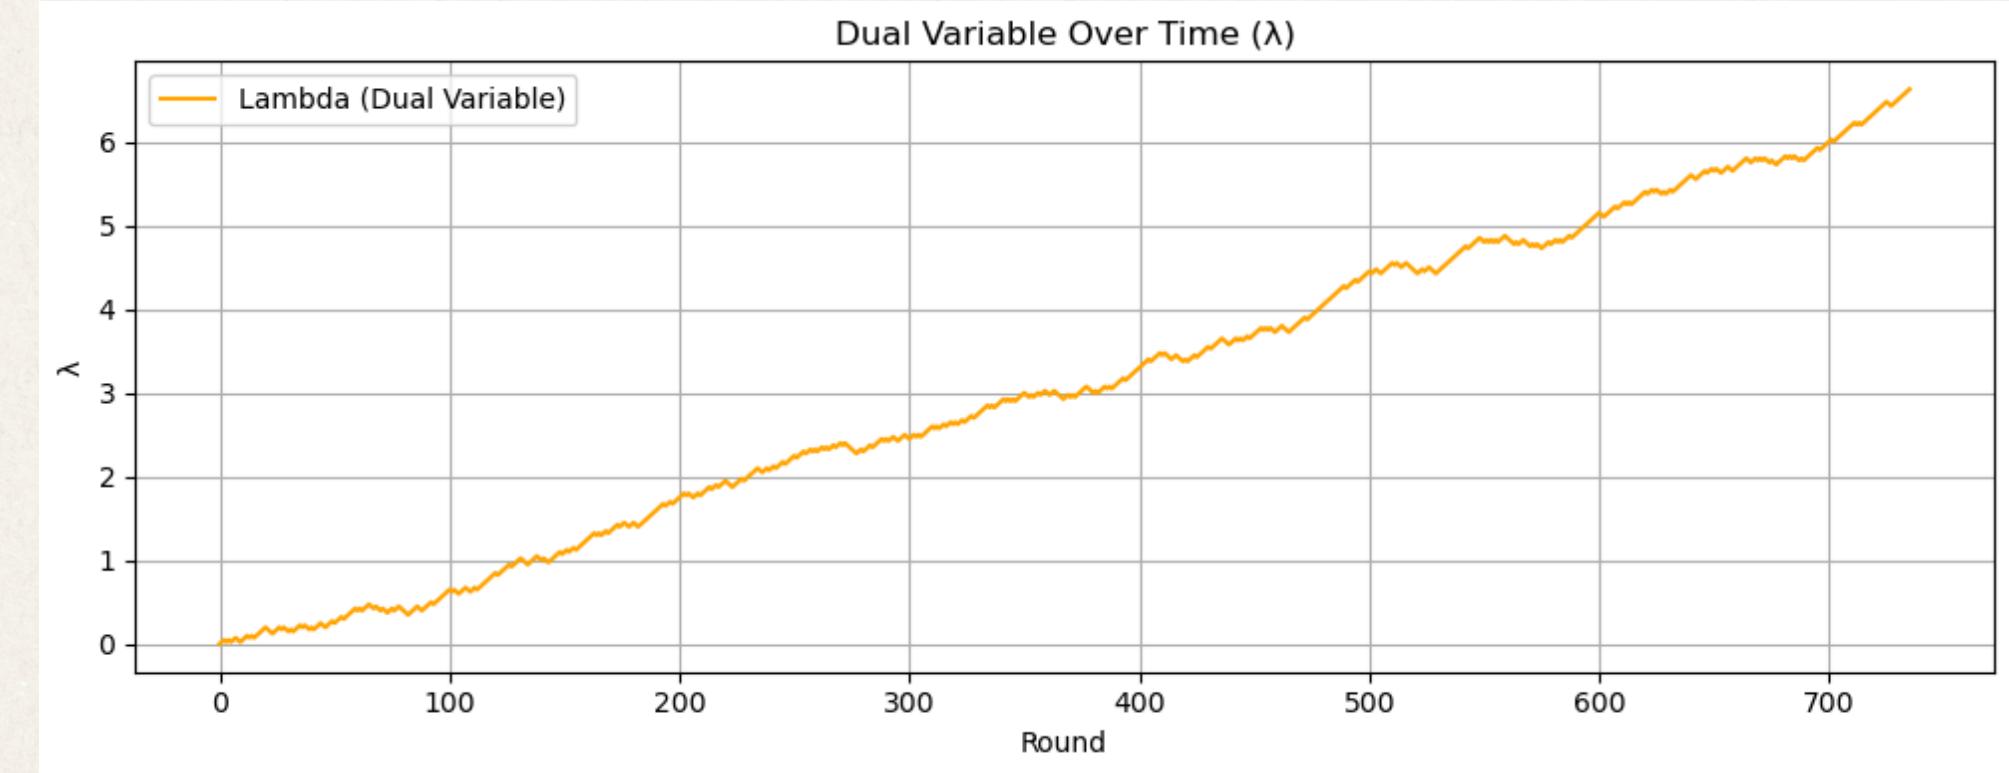
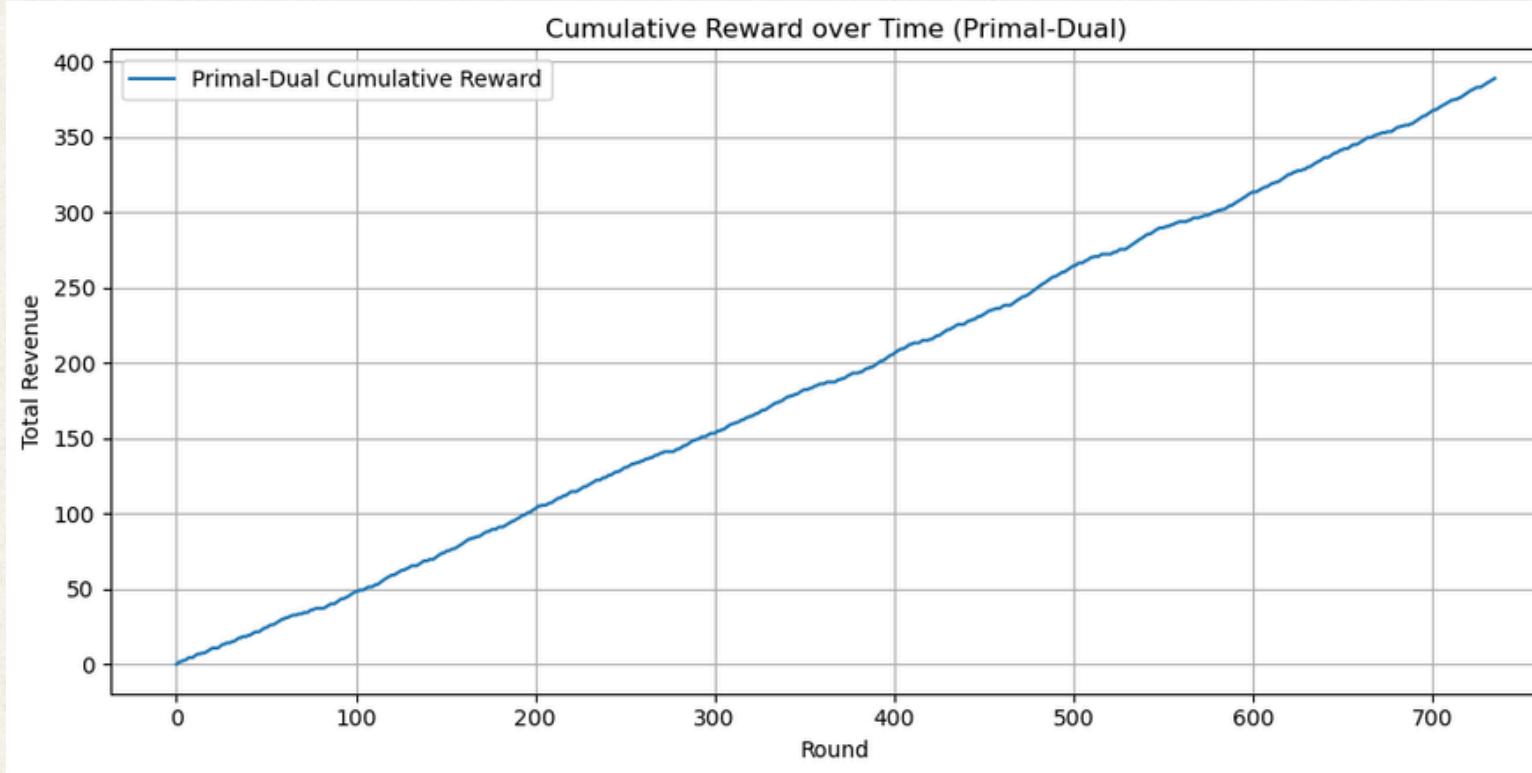
Observations:

- Primal-Dual adapts to non-stationary shifts.
- λ stabilizes over time, controlling inventory usage.

Settings:

- Budget $B = 500$
- Horizon $T = 1000$
- Prices: $\{0.1, 0.2, 0.3, 0.4, 0.5, 0.6, 0.7, 0.8, 0.9, 1.0\}$

Requirement 3: Best-of-Both-Worlds and single product



Requirement 4: Best-of-Both-Worlds and multiple products

Environment:

- Highly non-stationary valuation distribution.
- Joint distribution over valuations.

Settings:

- Budget $B = 8000$
- Horizon $T = 2000$
- Prices: $\{0.1, 0.23, 0.36, 0.5, 0.63, 0.76, 0.9\}$

Algorithms implemented:

- Primal-Dual method with inventory constraint.
- Separate regret minimiser per product.

Evaluation:

- Cumulative Regret Comparison.
- Cumulative Revenue Comparison.

Requirement 5: Slightly non-stationary environments with multiple products

Environment:

- Highly non-stationary valuation distribution.
- Joint distribution over valuations.

Settings:

- Budget $B = 8000$
- Horizon $T = 2000$
- Prices: $\{0.1, 0.23, 0.36, 0.5, 0.63, 0.76, 0.9\}$

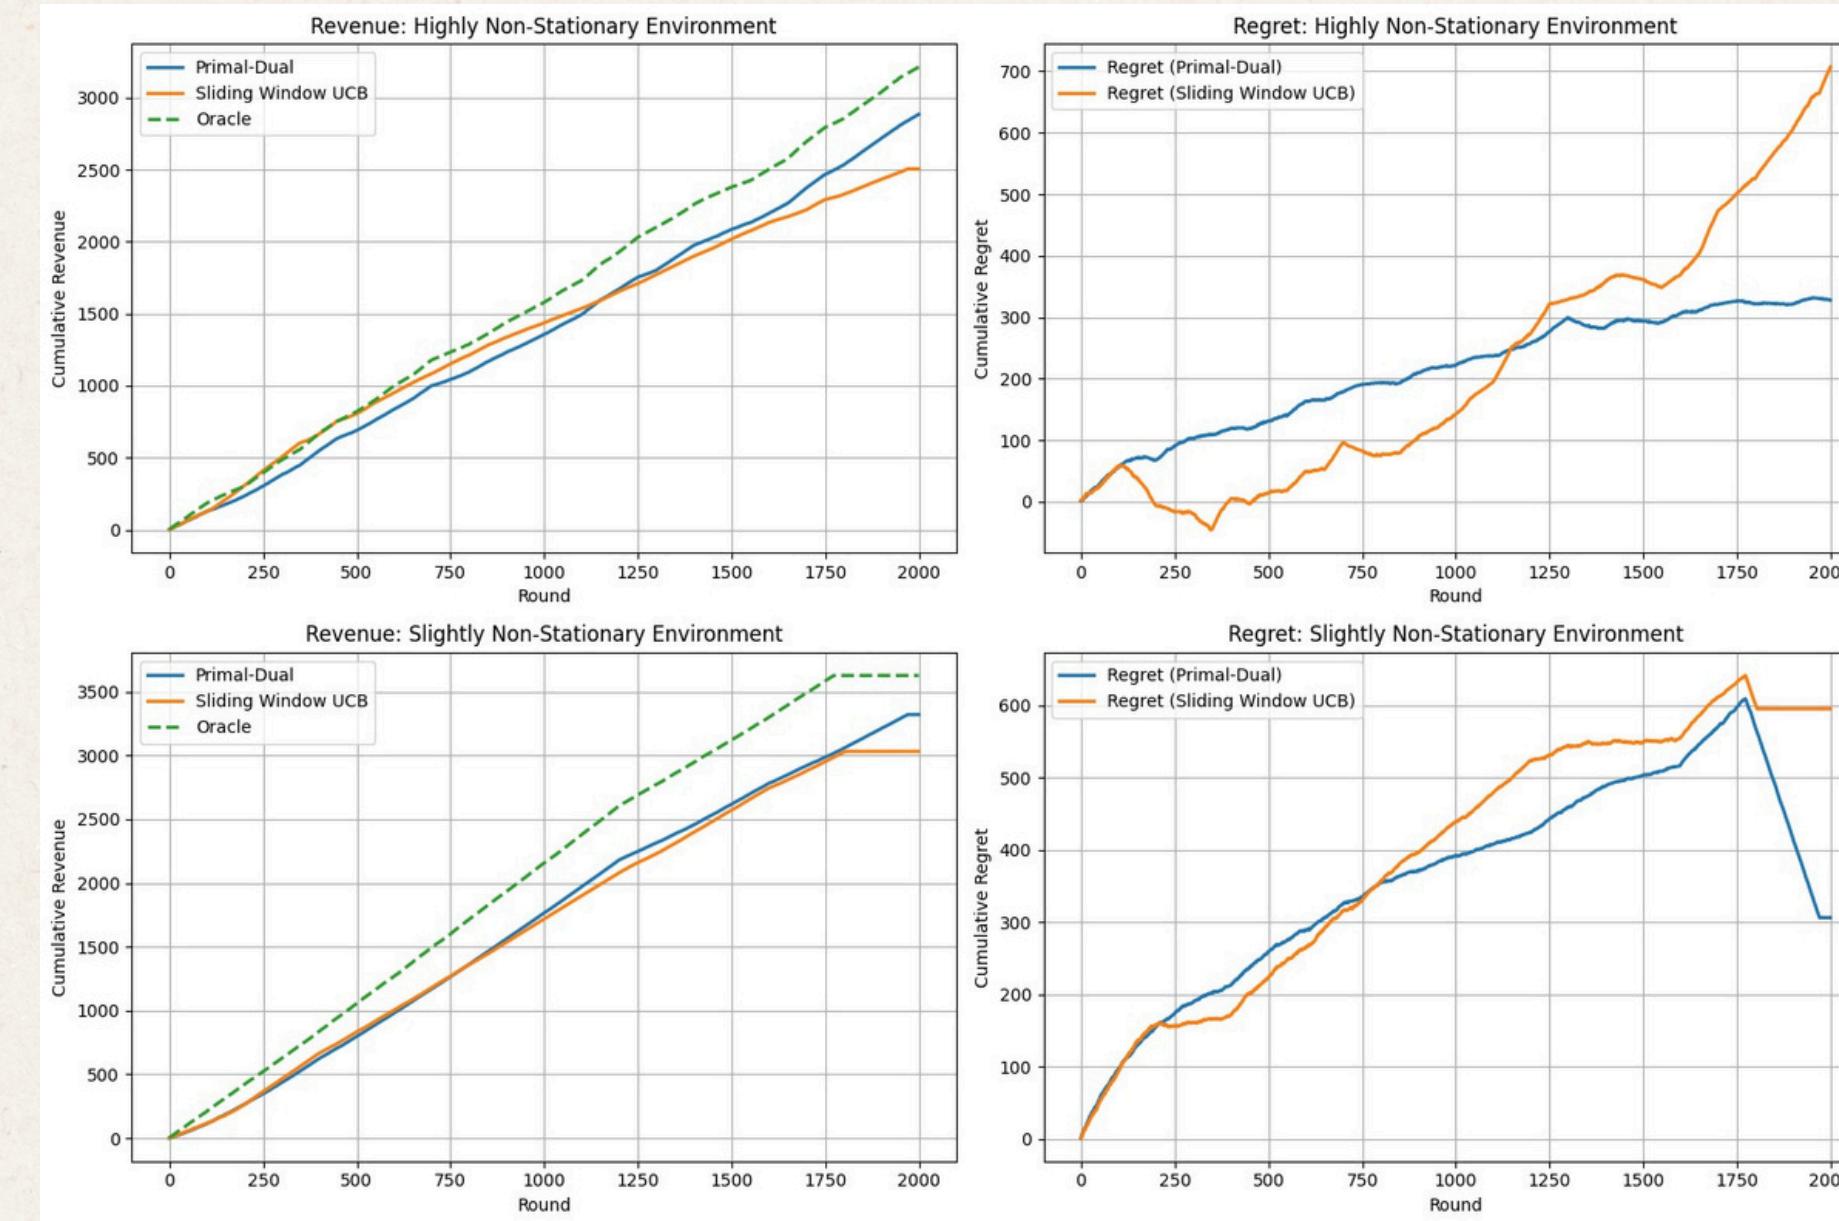
Algorithms implemented:

- Combinatorial-UCB with sliding window.

Evaluation:

- Cumulative Regret Comparison.
- Cumulative Revenue Comparison.

Requirement 4/5 results:



**Thank you for your
attention!**
